

# UNIX System Calls

# System Process

- The only active entities in a UNIX system are the processes.
- UNIX processes are very similar to the classical sequential processes.
- Each process runs a single program and initially has a single thread of control i.e. it has one program counter, which keeps track of the next instruction to be executed.
- UNIX is a multiprogramming system i.e. multiple, independent processes may be running simultaneously.

# System Process ...

- Each user may have several active processes at once, so on a large system, there may be hundreds or even thousands of processes running.
- In single-user workstations, even when the user is absent, dozens of background processes, called daemons (i.e. self employed evil spirit), are running.
- These daemons are started automatically when the system is booted.

# System Process ...

- “Cron Daemon” is a typical daemon.
- It wakes up once a minute to check if there is any work for it to do.
- If so, it does the work.
- Then it goes back to sleep until it is time for the next check.
- This daemon is needed because it is possible in UNIX to schedule activities minutes, hours, days, or even months in the future.

# System Process ...

- For example, suppose a user has an appointment with dentist at say 5:00 pm on 11.01.2000, he can make an entry in the cron daemon's database telling the daemon to beep at him at, say, 4:30 pm on 11.01.2000.
- When the appointed day and time arrives, the cron daemon sees that it has work to do, and starts up the beeping program as a new process.

# System Process ...

- Other daemons handle incoming and outgoing electronic mail, manage the line printer queue, check if there are enough free pages in memory, and so forth.
- Daemons are straightforward to implement in UNIX because each one is a separate process, independent of all other processes.

# System Process ...

- For example, suppose a user has a dentist appointment at 3 o'clock next Tuesday.
- He can make an entry in the cron daemon's database telling the daemon to beep at him at, say, 2:30.
- When the appointed day and time arrives, the cron daemon sees that it has work to do, and starts up the beeping program as a new process.

# System Process ...

- The cron daemon is also used to start up periodic activities, such as making daily disk backups at say 4 A.M., or reminding forgetful users every year on any particular date or month i.e. say December.



# System Process ...

- Processes are created in UNIX in an especially simple manner.
- The fork system call creates an exact copy of the original process.
- The forking process is called the parent process.
- The new process is called the child process.
- The parent and child each have their own, private memory images.

# System Process ...

- If the parent subsequently changes any of its variables, the changes are not visible to the child, and vice versa.
- Open files are shared between parent and child. That is, if a certain file was open in the parent before the fork, it will continue to be open in both the parent and the child afterward.
- Changes made to the file by either one will be visible to the other.
- This behavior is only reasonable, because these changes are also visible to any unrelated process that opens the file as well.

# System Process ...

- The memory images, variables, registers, and everything else are identical in the parent and child leads to a small difficulty:
- How do the processes know which one should run the parent code and which one should run the child code?
- The secret is that the fork system call returns a 0 to the child and a nonzero value, the child's PID (Process Identifier) to the parent. Processes are named by their PIDs.
- When a process is created, the parent is given the child's PID.

# System Process ...

- If the child wants to know its own PID, there is a system call, getpid, that provides it.
- PIDs are used in different of ways.
- For example, when a child terminates, the parent is given the PID of the child that just finished.
- This can be important because a parent may have many children.

# System Process ...

- Since children may also have children, an original process can build up an entire tree of children, grandchildren, and further descendants.
- Processes in UNIX can communicate with each other using a form of message passing.
- It is possible to create a channel between two processes into which one process can write a stream of bytes for the other to read.
- These channels are called pipes.

# System Process ...

- Synchronization is possible because when a process tries to read from an empty pipe it is blocked until data are available.
- Shell pipelines are implemented with pipes. When the shell sees a line like “sort <f | head”, it creates two processes, sort and head, and sets up a pipe between them in such a way that standard output of sort is connected to that of head standard input.

# System Process ...

- In this way, all the data that sort writes go directly to head, instead of going to a file. If the pipe fills up, the system stops running sort until head has removed some data from the pipe.
- Processes can also communicate through software interrupts.
- A process can send what is called a signal to another process.

# System Process ...

- Processes can tell the system what they want to happen when a signal arrives. The choices are
  - to ignore it
  - to catch it
  - to let the signal kill the process (the default for most signals)



# System Process ...

- If a process elects to catch signals sent to it, it must specify a signal handling procedure.
- When a signal arrives, control will abruptly switch to the handler.
- When the handler is finished and returns, control goes back to where it came from, analogous to hardware I/O interrupts.
- A process can only send signals to members of its process group, which consists of its parent (and further ancestors), siblings, and children (and further descendants).

# System Process ...

- A process may also send a signal to all members of its process group with a single system call.
- Signals are also used for other purposes.
- For example, if a process is doing floating-point arithmetic, and inadvertently divides by 0, it gets a SIGFPE (Floating-Point Exception) signal.
- The signals that are required by POSIX [Portable Operating System Interface] are listed given below ...

# System Process ...

| Signal  | Cause   |
|---------|---|
| SIGABRT | Sent to abort a process and force a core dump     |
| SIGALRM | Alarm clock has gone off                          |
| SIGFPE  | Floating point error has occurred                 |
| SIGHUP  | Phone line the process was using has hung up      |
| SIGILL  | User has hit the DEL key to interrupt the process |
| SIGQUIT | User has hit the key requesting the core dump     |

# System Process ...

| Signal  | Cause   |
|---------|---|
| SIGKILL | Sent to kill a process                              |
| SGPIPE  | The process written to a pipe which has no header   |
| SIGSEGV | Process has reference an invalid memory address     |
| SIGTERM | Used to request that a process terminate gracefully |
| SIGUSR1 | Available for application defined purpose           |
| SIGUSR2 | Available for application defined purpose           |

# System Process ...

| System Call                       | Decription                                    |
|-----------------------------------|---|
| pid=fork()                        | Create a child process identical to parent    |
| pid=waitpid(pid, staloc, opts)    | Wait for child to terminate                   |
| s= execve(name, argv, envp)       | Replace the core image of process             |
| exit(status)                      | Terminate process execution and return status |
| s = sigaction(sig, &act, &oldact) | Define action to take on signals              |
| s = sigreturn(&context)           | Return from a signal                          |
| s = sigprocmask(how, &set, &old)  | Examine or change the signal mask             |

# System Process ...

| System Call                            | Description                                     |
|--|---|
| <code>s = sigpending(set)</code>       | Get the set of blocked signals                  |
| <code>s = sigsuspend(sigmask)</code>   | Replace the signal mask and suspend the process |
| <code>s = kill(pid, sig)</code>        | Send a signal to a process                      |
| <code>residual = alarm(seconds)</code> | Set the alarm clock                             |
| <code>s = pause( )</code>              | Suspend the caller until the next signal        |

# Unix System Calls

- “**System Calls**” are assembly language instructions.
- The “**System Call**” provide interface between a process and an operating system.
- “**System calls**” are usually made when a process in user mode requires access to a resource.
- The “**System Call**” requests the kernel to provide the resource.

# Types of Unix System Calls

- Process Control
- File Management
- Device Management
- Information Maintenance
- Communication



# System Calls Windows v/s UNIX

| Types of System Calls | Windows  | Linux                                  |
|-----------------------|--|--|
| Process Control       | CreateProcess()<br>ExitProcess()<br>WaitForSingleObject()  | fork()<br>exit()<br>wait()             |
| File Management       | CreateFile()<br>ReadFile()<br>WriteFile()<br>CloseHandle() | open()<br>read()<br>write()<br>close() |

# System Calls Windows v/s UNIX

| Types of System Calls   | Windows  | Linux                          |
|-------------------------|--|--------------------------------|
| Device Management       | SetConsoleMode()<br>ReadConsole()<br>WriteConsole()    | ioctl()<br>read()<br>write()   |
| Information Maintenance | GetCurrentProcessID()<br>SetTimer()<br>Sleep()         | getpid()<br>alarm()<br>sleep() |
| Communication           | CreatePipe()<br>CreateFileMapping()<br>MapViewOfFile() | pipe()<br>shmget()<br>mmap()   |

# Types of Unix System Calls: Process Control

- These system calls deal with processes like
  - Process Creation
  - Process Termination
  - ...

# Types of Unix System Calls: File Management

- These system calls are responsible for file manipulation like
  - Creating File
  - Reading File
  - Writing into File
  - ...

# Types of Unix System Calls: Device Management

- These system calls are responsible for device manipulation like reading from device buffers, writing into device buffers etc.
  - Reading from Device Buffer
  - Writing into Device Buffer
  - ...

# Types of Unix System Calls: Information Maintenance

- These system calls handle information and its transfer between the operating system and the user program.

# Types of Unix System Calls: Communication

- These system calls are useful for interprocess communication. They also deal with creating and deleting a communication connection.

# Example of Unix System Call

- “open()”
- “read()”
- “creat()”
- “write()”
- “lseek()”
- “close()”
- “stat()”



# Example of Unix System Call

- “fstat()”
- “dup()”
- “link()”
- “access()”
- “chmod()”
- “chown()”
- “unmask()”

# Example of Unix System Call

- “ulink()”
- “exec()”
- “fork()”
- “wait()”

# Unix System Calls: “exec()”

- This system call runs an executable file in the context of an already running process.
- It replaces the previous executable file.
- This is known as an overlay.
- The original process identifier remains since a new process is not created but data, heap, stack etc. of the process are replaced by the new process.

# Unix System Calls: “exit()”

- The `exit()` system call is used by a program to terminate its execution.
- In a multithreaded environment, this means that the thread execution is complete.
- The operating system reclaims resources that were used by the process after the `exit()` system call.

# Unix System Calls: “fork()”

- Processes use the fork() system call to create processes that are a copy of themselves.
- This is one of the major methods of process creation in operating systems.
- When a parent process creates a child process and the execution of the parent process is suspended until the child process executes.
- When the child process completes execution, the control is returned back to the parent process.

# Unix System Calls: “kill()”

- The “kill() system call” is used by the operating system to send a termination signal to a process that urges the process to exit.
- The “kill system call” does not necessary mean killing the process and can have different meaning(s).

# Unix System Calls: “wait()”

- In some systems, a process may wait for another process to complete its execution.
- This happens when a parent process creates a child process and the execution of the parent process is suspended until the child process executes.
- The suspending of the parent process occurs with a wait() system call.
- When the child process completes execution, the control is returned back to the parent process.